Exam 1 review

Questions

> **i** Note
>
> Suggested answers can be found here, but resist the urge to peek before you go through it yourself.

Blizzard salaries

In 2020, employees of Blizzard Entertainment circulated a spreadsheet to anonymously share salaries and recent pay increases amidst rising tension in the video game industry over wage disparities and executive compensation. (Source: Blizzard Workers Share Salaries in Revolt Over Pay)

The name of the data frame used for this analysis is `blizzard_salary` and the variables are:

- `percent_incr`: Raise given in July 2020, as percent increase with values ranging from 1 (1% increase) to 21.5 (21.5% increase)
- `salary_type`: Type of salary, with levels `Hourly` and `Salaried`
- `annual_salary`: Annual salary, in USD, with values ranging from $50,939 to $216,856.
- `performance_rating`: Most recent review performance rating, with levels `Poor`, `Successful`, `High`, and `Top`. The `Poor` level is the lowest rating and the `Top` level is the highest rating.

The first ten rows of `blizzard_salary` are shown below:

```
# A tibble: 409 × 4
   percent_incr salary_type annual_salary performance_rating
          <dbl> <chr>               <dbl> <chr>
 1          1   Salaried                1  High
 2          1   Salaried                1  Successful
 3          1   Salaried                1  High
 4          1   Hourly              33987. Successful
 5         NA   Hourly              34798. High
 6         NA   Hourly              35360  <NA>
 7         NA   Hourly              37440  <NA>
 8          0   Hourly              37814. <NA>
 9          4   Hourly              41101. Top
10          1.2 Hourly              42328  <NA>
# ℹ 399 more rows
```

Question 1

Which of the following is **correct**? Choose all that apply.

a. The `blizzard_salary` dataset has 399 rows.

b. The `blizzard_salary` dataset has 4 columns.

c. Each row represents a Blizzard Entertainment worker who filled out the spreadsheet.

d. The `percent_incr` variable is numerical and discrete.

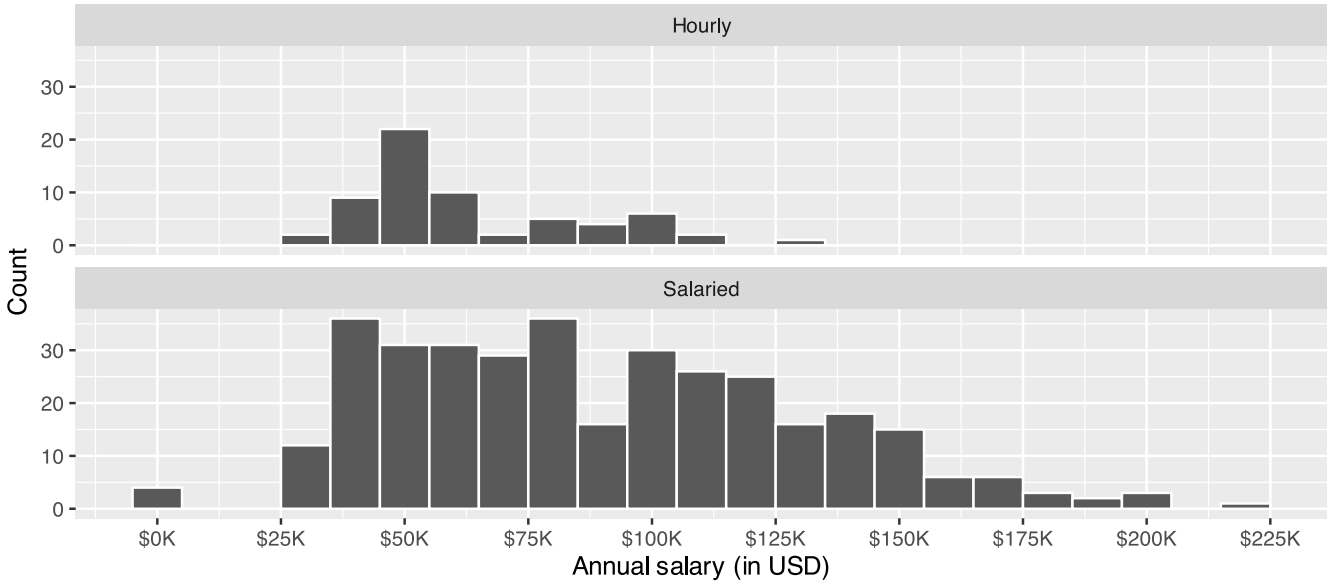
e. The `salary_type` variable is numerical.

f. The `annual_salary` variable is numerical.

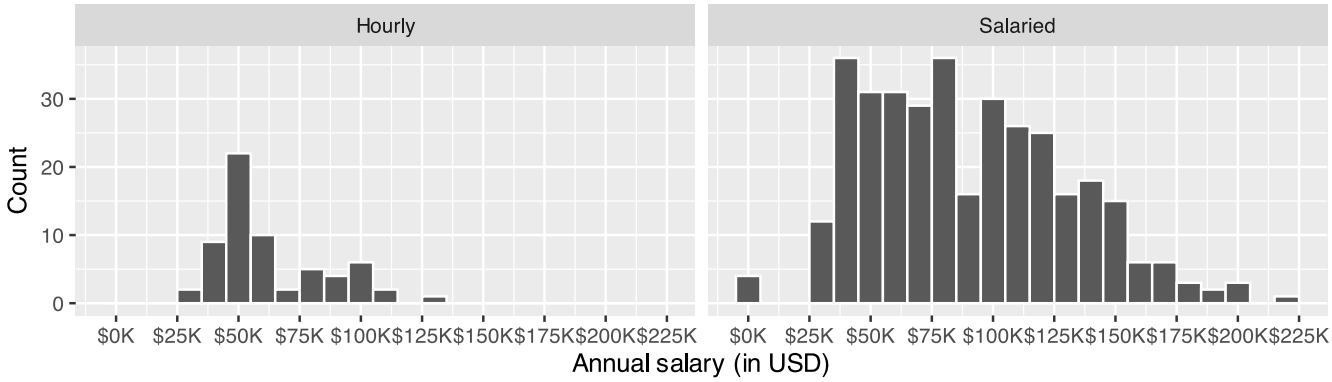
g. The `performance_rating` variable is categorical and ordinal.

Question 2

Figure 1a and Figure 1b show the distributions of annual salaries of hourly and salaried workers. The two figures show the same data, with the facets organized across rows and across columns. Which of the two figures is better for comparing the median annual salaries of hourly and salaried workers. Explain your reasoning.



(a) Option 1



(b) Option 2

Figure 1: Distribution of annual salaries of Blizzard employees

Question 3

Suppose your teammate wrote the following code as part of their analysis of the data.

```
blizzard_salary |>
  group_by(salary_type) |>
  summarize(
    mean_annual_salary = mean(annual_salary),
    median_annual_salary = median(annual_salary)
  )
```

They then printed out the results shown below. Unfortunately one of the numbers got erased from the printout. It's indicated with _____ below.

```
# A tibble: 2 × 3
  salary_type mean_annual_salary median_annual_salary
  <chr>                    <dbl>                <dbl>
1 Hourly                  63003.               54246.
2 Salaried                90183.               _____
```

Which of the following is the best estimate for that erased value?

a. 30,000

b. 50,000

c. 80,000

d. 100,000

Question 4

Which distribution of annual salaries has a higher standard deviation?

a. Hourly workers

b. Salaried workers

c. Roughly the same

Question 5

Which of the following alternate plots would also be useful for visualizing the distributions of annual salaries of hourly and salaried workers? Choose all that apply.

a. Box plots

b. Density plots

c. Pie charts

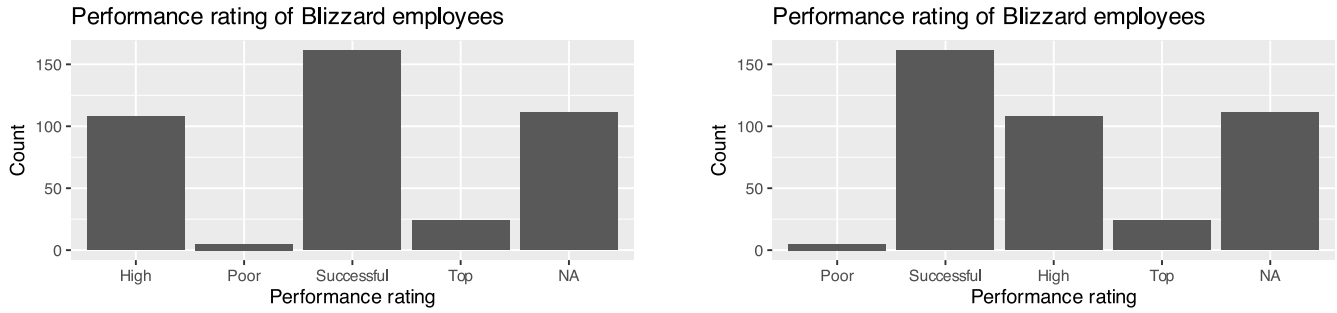
d. Waffle charts

e. Histograms

f. Scatterplots

Questions 6 and 7

Suppose you made the bar plot shown in Figure 2a to visualize the distribution of `performance_rating` and your teammate made the bar plot shown in Figure 2b.



(a) Option 1 (b) Option 2

Figure 2: Distribution of performance rating

You made your bar plot without transforming the data in any way, while your friend did first transform the data with code like the following:

```
blizzard_salary <- blizzard_salary |>
  _(1)_(performance_rating = fct_relevel(performance_rating, _(2)_))
```

Question 6: What goes in the blank (1)?

a. `arrange()`

b. `filter()`

c. `mutate()`

d. `summarize()`

Question 7: What goes in the blank (2)?

a. `"Poor", "Successful", "High", "Top"`

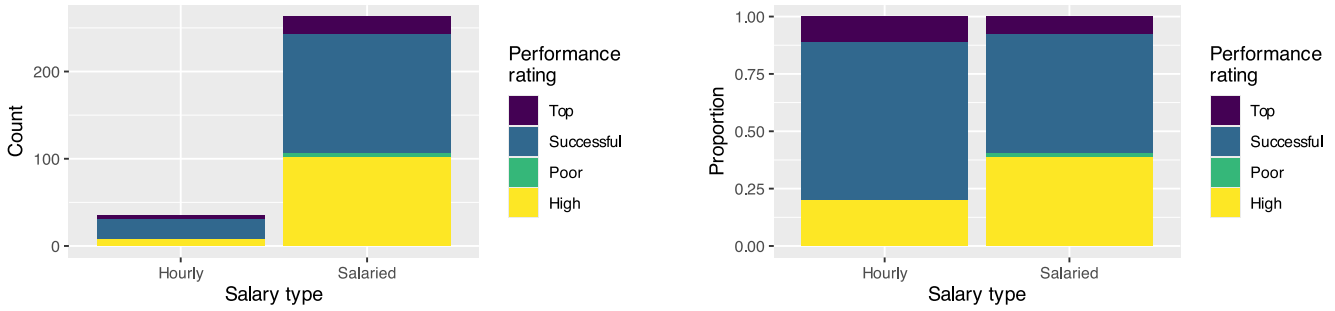
b. `"Successful", "High", "Top"`

c. `"Top", "High", "Successful", "Poor"`

d. `Poor, Successful, High, Top`

Questions 8 - 10

Finally, another teammate creates the following two plots.



(a) Option 1 (b) Option 2

Figure 3: Distribution of salary type by performance rating

Question 8: Your teammate asks you for help deciding which one to use in the final report for visualizing the relationship between performance rating and salary type. In 1-3 sentences, can you help them make a decision, justify your choice, and write the narrative that should go with the plot?

Question 9: A friend with a keen eye points out that the number of observations in Figure 3a seems lower than the total number of observations in `blizzard_salary`. What might be going on here? Explain your reasoning.

Question 10: Below are the proportions of performance ratings for hourly and salaried workers. Place these values in the corresponding segments in Figure 3b.

```
# A tibble: 4 × 3
  performance_rating Hourly Salaried
  <chr>               <dbl>    <dbl>
1 High                0.2     0.384
2 Successful          0.686   0.521
3 Top                 0.114   0.0760
4 Poor                0       0.0190
```

Questions 11 and 12

The table below shows the distribution of `salary_type` and `performance_rating`.

```
# A tibble: 2 × 6
  salary_type  Poor Successful  High   Top  `NA`
  <chr>       <int>      <int> <int> <int> <int>
1 Hourly         NA         24     7     4    28
2 Salaried        5        137   101    20    83
```

The pipeline below produces a data frame with a fewer number of rows than `blizzard_salary`.

```
blizzard_salary |>
  filter(salary_type _(1)_ "Hourly" _(2)_ performance_rating == "Poor") |>
  _(3)_(annual_salary)
```

```
# A tibble: 5 × 4
  percent_incr salary_type annual_salary performance_rating
         <dbl> <chr>              <dbl> <chr>
1            0 Salaried           80000 Poor
2            3 Salaried           83000 Poor
3            0 Salaried          116000 Poor
4            0 Salaried          135219 Poor
5            0 Salaried          147500 Poor
```

Question 11: Which of the following goes in blanks (1) and (2)?

	(1)	(2)
a.	`!=`	`|`
b.	`==`	`&`
c.	`!=`	`&`
d.	`==`	`|`

Question 12: Which function or functions go into blank (3)?

a. `arrange()`

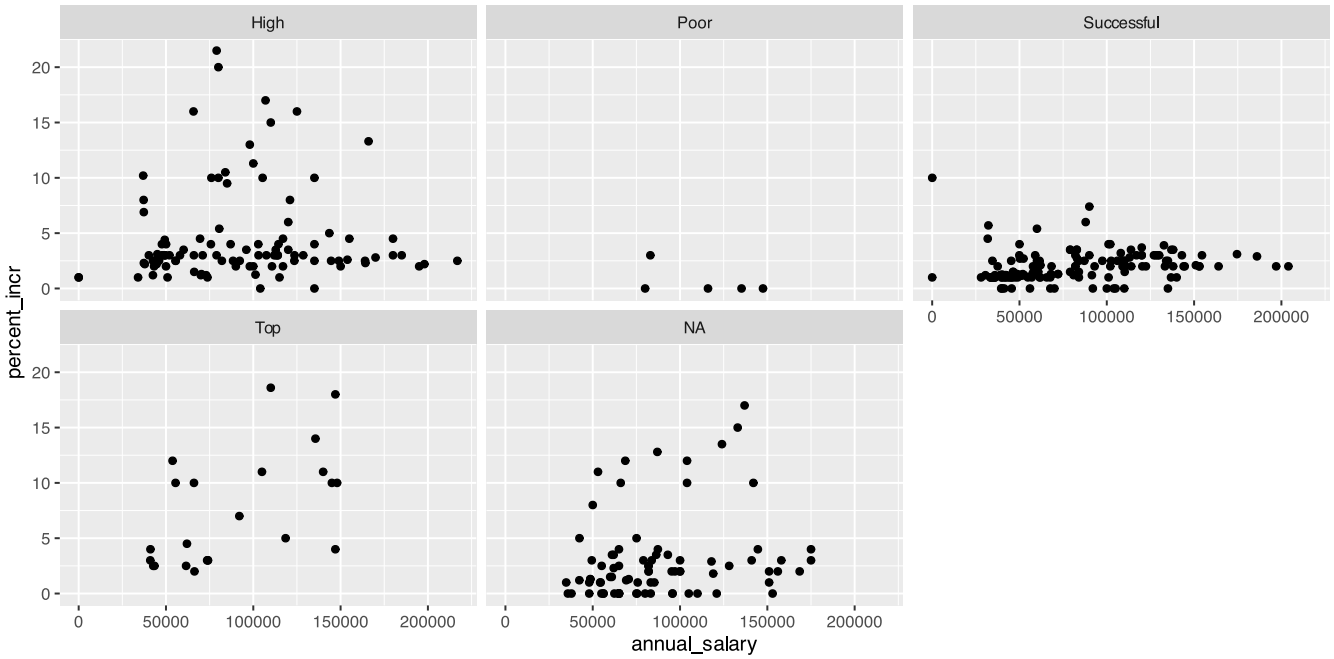
b. `mutate()`

c. `order()`

d. `sort()`

Question 13

You're reviewing another team's work and they made the following visualization:



And they wrote the following interpretation for the relationship between annual salary and percent increase for Top performers:

> The relationship is positive, having a higher salary results in a higher percent increase. There is one clear outlier.

Which of the following is/are the most accurate and helpful) peer review note for this interpretation. Choose all that apply.

a. The interpretation is complete and perfect, no changes needed!

b. The interpretation doesn't mention the direction of the relationship.

c. The interpretation doesn't mention the form of the relationship, which is linear.

d. The interpretation doesn't mention the strength of the relationship, which is somewhat strong.

e. There isn't a clear outlier in the plot. If any points stand out as potential outliers, more guidance should be given to the reader to identify them (e.g., salary and/or percent increase amount).

f. The interpretation is causal – we don't know if the cause of the high percent increase is higher annual salary based on observational data. The causal direction might be the other way around, or there may be other factors contributing to the apparent relationship.

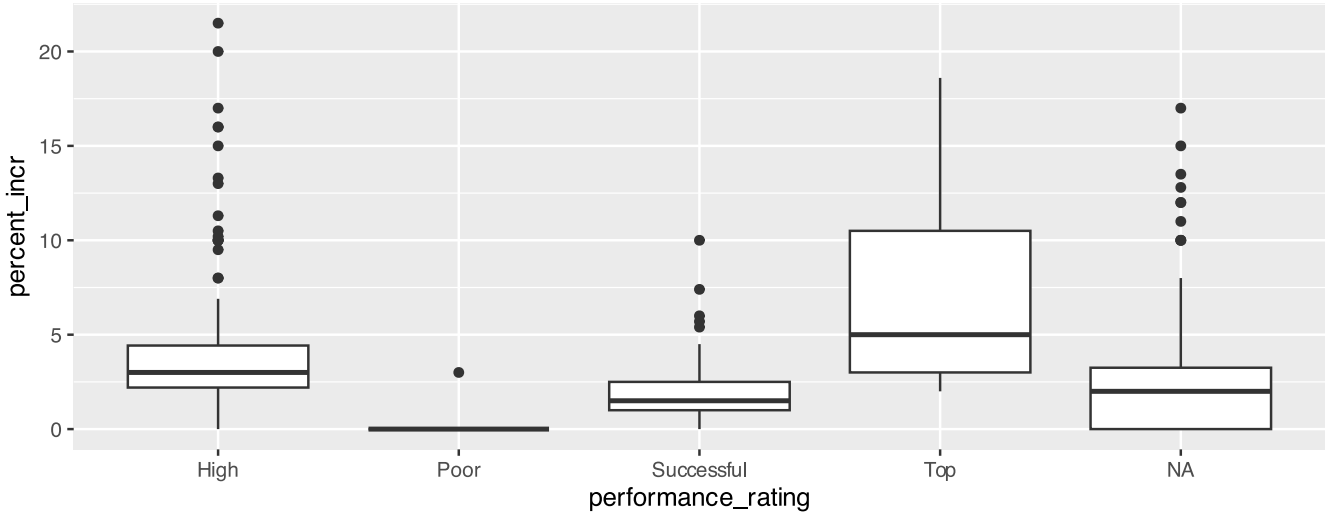
Question 14

Below is some code and its output.

```
```{r}
# label=plot blizzard

ggplot(blizzard_salary,aes(x=performance_rating,y=percent_incr))+geom_boxplot()
labs(x="Performance rating", y = "Percent increase")
```
```

```
Warning: Removed 39 rows containing non-finite outside the scale range
(`stat_boxplot()`).
```




```
<ggplot2::labels> List of 2
 $ x: chr "Performance rating"
 $ y: chr "Percent increase"
```

Part 1: List at least 5 things that should be fixed or improved in the code.

Part 2: What is the cause of the warning and what does it mean?

Question 15

You're working on a data analysis on salaries of Blizzard employees in a Quarto document in a project version controlled by Git. You create a plot and write up a paragraph describing any patterns in it. Then, your teammate says "render, commit, and push".

Part 1: What do they mean by each of these three steps. In 1-2 sentences for each, explain in your own words what they mean.

1. Render:

2. Commit:

3. Push:

Part 2: Your teammate is getting impatient and they interrupt you after you rendered and committed and say "I still can't see your changes in our shared GitHub repo when I look at it in my web browser." Which of the following answers is the most accurate?

a. I rendered my document, you should be seeing my changes on GitHub when you look at it in your web browser.

b. I committed my changes, you should be seeing my changes on GitHub when you look at it in your web browser.

c. I didn't yet push my changes, it's expected that you are not seeing them on GitHub when you look at it in your web browser. Wait until I push, and check again.

d. You need to pull to see my changes on GitHub in the web browser.

Penguins

The `penguins` data set includes measurements for penguin species, including: flipper length, body mass, bill dimensions, and sex. The following table summarizes information on which species of penguins (Adelie, Gentoo, and Chinstrap) live on which islands (Biscoe, Dream, or Torgersen).

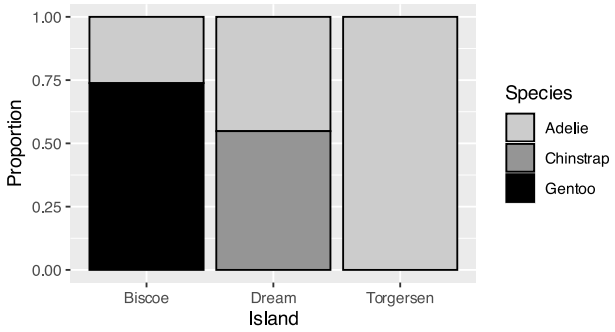
Island	Adelie	Gentoo	Chinstrap	Total
Biscoe	44	124	0	168
Dream	56	0	68	124
Torgersen	52	0	0	52
Total	152	124	68	344

Question 16

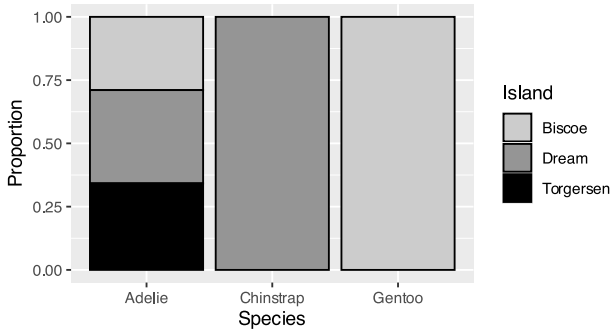
Which of the following plots is the result of the following code?

```
ggplot(penguins, aes(x = island, fill = species)) +
  geom_bar()
```

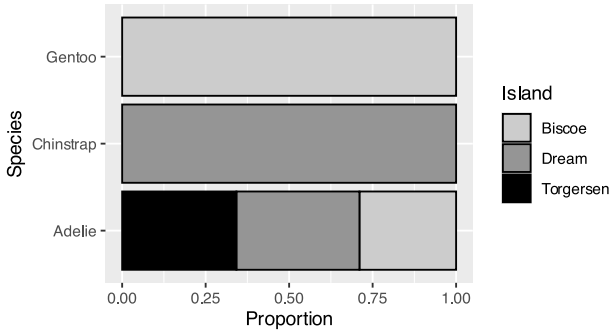
(a)



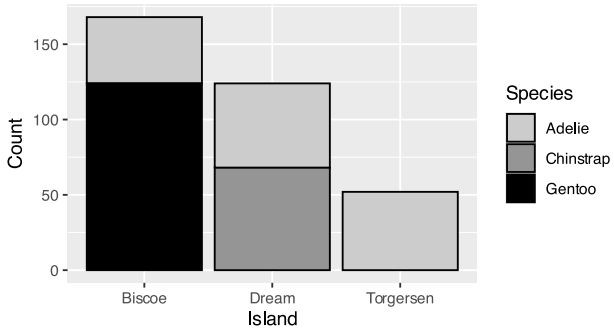
(b)



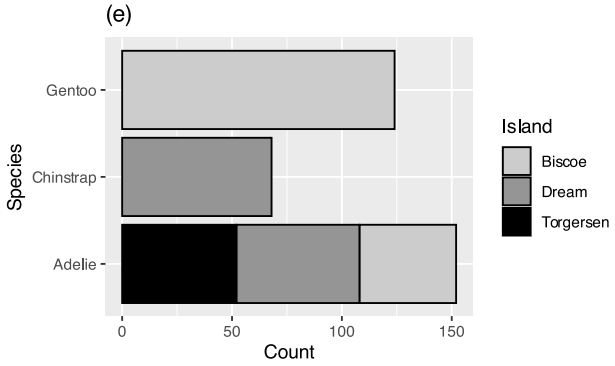
(c)



(d)



(e)



NYC Flights

The `flights` dataset includes characteristics of all flights departing from New York City airports (JFK, LGA, EWR) in 2013. Below is a peek at the first ten rows of the `flights` data.

```
# A tibble: 336,776 × 19
    year month   day arr_delay carrier dep_time sched_dep_time dep_delay
   <int> <int> <int>     <dbl> <chr>      <int>          <int>     <dbl>
 1  2013     1     1        11 UA           517            515         2
 2  2013     1     1        20 UA           533            529         4
 3  2013     1     1        33 AA           542            540         2
 4  2013     1     1       -18 B6           544            545        -1
 5  2013     1     1       -25 DL           554            600        -6
 6  2013     1     1        12 UA           554            558        -4
 7  2013     1     1        19 B6           555            600        -5
 8  2013     1     1       -14 EV           557            600        -3
 9  2013     1     1        -8 B6           557            600        -3
10  2013     1     1         8 AA           558            600        -2
# ℹ 336,766 more rows
# ℹ 11 more variables: arr_time <int>, sched_arr_time <int>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>
```

Question 17

Based on this output, which of the following must be **true** about the `flights` data frame? **Select all that are true.**

a. The `flights` data frame is a `tibble`.

b. The `flights` data frame has 10 rows.

c. The `flights` data frame has 8 columns.

d. The `carrier` variable in the `flights` data frame is a character variable.

e. There are no missing data in the `flights` data frame.

Question 18

Which of the following pipelines produce(s) the output shown below? **Select all that apply.**

```
# A tibble: 336,776 × 5
   arr_delay carrier  year month   day
       <dbl> <chr>   <int> <int> <int>
 1      1272 HA       2013     1     9
 2      1127 MQ       2013     6    15
 3      1109 MQ       2013     1    10
 4      1007 AA       2013     9    20
 5       989 MQ       2013     7    22
 6       931 DL       2013     4    10
 7       915 DL       2013     3    17
 8       895 DL       2013     7    22
 9       878 AA       2013    12     5
10       875 MQ       2013     5     3
# ℹ 336,766 more rows
```

a.

```
flights |>
  select(arr_delay, carrier, year, month, day) |>
  arrange(desc(arr_delay))
```

b.

```
flights |>
  select(arr_delay, carrier, year, month, day) |>
  arrange(arr_delay)
```

c.

```
flights |>
  select(arr_delay, carrier, year, month, day) |>
  arrange(year)
```

d.

```
flights |>
  arrange(desc(arr_delay)) |>
  select(arr_delay, carrier, year, month, day)
```

e.

```
flights |>
  arrange(desc(arr_delay)) |>
  select(day, month, year, arr_delay, carrier)
```

Countries and populations

We have a small dataset of six countries and their populations:

```
# A tibble: 6 × 2
  country       population
  <chr>              <dbl>
1 Curacao              150
2 Ecuador            18001
3 Iraq               44496.
4 New Zealand         5124.
5 Palau                 18.0
6 United States    333288.
```

And another small dataset of five countries and the continent they're in:

```
# A tibble: 5 × 3
  entity      code  continent
  <chr>       <chr> <chr>
1 Angola      AGO   Africa
2 Curacao     CUW   North America
3 Ecuador     ECU   South America
4 Iraq        IRQ   Asia
5 New Zealand NZL   Oceania
```

You join the two datasets with the following:

```
population |>
  left_join(continents, by = join_by(country == entity))
```

Question 19

How many rows will the resulting data frame have?

a. 4
b. 5
c. 6
d. 7
e. 8

Question 20

What will be the columns of the resulting data frame?

a. `country`, `population`

b. `country`, `population`, `code`, `continent`

c. `entity`, `code`, `continent`

d. `entity`, `population`, `code`, `continent`

e. `country`, `entity`, `population`, `code`, `continent`

Duke Forest houses

The `duke_forest` dataset includes information on prices and various other features (number of bedrooms, bathrooms, area, year built, type of cooling, type of heating, etc.) of houses in the Duke Forest neighborhood of Durham, NC.

```
Rows: 98
Columns: 13
$ address    <chr> "1 Learned Pl, Durham, NC 27705", "1616 Pinecrest Rd, Durha…
$ price      <dbl> 1520000, 1030000, 420000, 680000, 428500, 456000, 1270000, …
$ bed        <dbl> 3, 5, 2, 4, 4, 3, 5, 4, 4, 3, 4, 4, 3, 5, 4, 5, 3, 4, 4, 3,…
$ bath       <dbl> 4.0, 4.0, 3.0, 3.0, 3.0, 3.0, 5.0, 3.0, 5.0, 2.0, 3.0, 3.0,…
$ area       <dbl> 6040, 4475, 1745, 2091, 1772, 1950, 3909, 2841, 3924, 2173,…
$ type       <chr> "Single Family", "Single Family", "Single Family", "Single …
$ year_built <dbl> 1972, 1969, 1959, 1961, 2020, 2014, 1968, 1973, 1972, 1964,…
$ heating    <chr> "Other, Gas", "Forced air, Gas", "Forced air, Gas", "Heat p…
$ cooling    <fct> central, central, central, central, central, central, centr…
$ parking    <chr> "0 spaces", "Carport, Covered", "Garage - Attached, Covered…
$ lot        <dbl> 0.97, 1.38, 0.51, 0.84, 0.16, 0.45, 0.94, 0.79, 0.53, 0.73,…
$ hoa        <chr> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,…
$ url        <chr> "https://www.zillow.com/homedetails/1-Learned-Pl-Durham-NC-…
```

The following summary table gives us some information about whether homes in this data set have garages and when they were built.

	Built earlier than 1950	Built in 1950 or later
Garage	5	33
No garage	3	57

The pipeline below produces a data frame with a fewer number of rows than `duke_forest`.

```
duke_forest |>
  filter(parking == "Garage" _(1)_ year_built _(2)_ 1950) |>
  select(parking, year_built, price, area) |>
  _(3)_(price_per_sqfeet = price / area)
```

```
# A tibble: 5 × 5
  parking year_built  price  area price_per_sqfeet
  <chr>        <dbl>  <dbl> <dbl>           <dbl>
1 Garage        1945 900000  2933            307.
2 Garage        1938 265000  1300            204.
3 Garage        1934 600000  2514            239.
4 Garage        1941 412500  1661            248.
5 Garage        1940 105000  1094             96.0
```

Question 21

Which of the following goes in blanks (1) and (2)?

	(1)	(2)
a.	`&`	`<`
b.	`|`	`<`
c.	`&`	`>=`
d.	`|`	`>=`
e.	`&`	`!=`

Question 22

Which function or functions go into blank (3)? **Select all that apply.**

a. `arrange()`

b. `mutate()`

c. `filter()`

d. `summarize()`

e. `slice()`

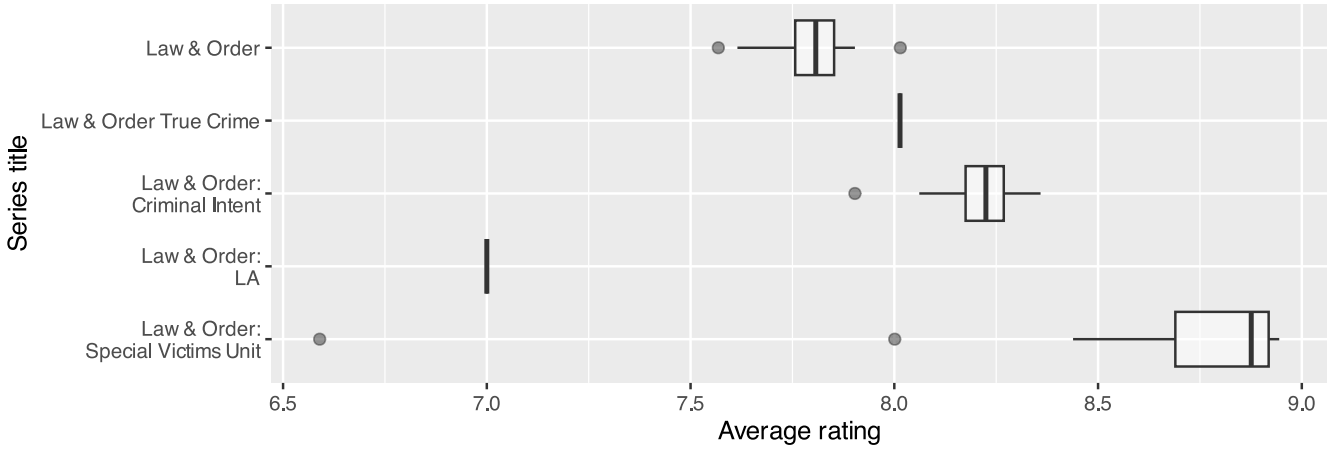
Law & Order

You've heard of the tidyverse, now let's visit the Law & Order-verse. Doink doink![1]

Law & Order is a police procedural and legal drama television series that has been running since the 1990s. The Law & Order franchise includes a number of series such as Law & Order, Law & Order: SVU, Law & Order: Criminal Intent, etc.

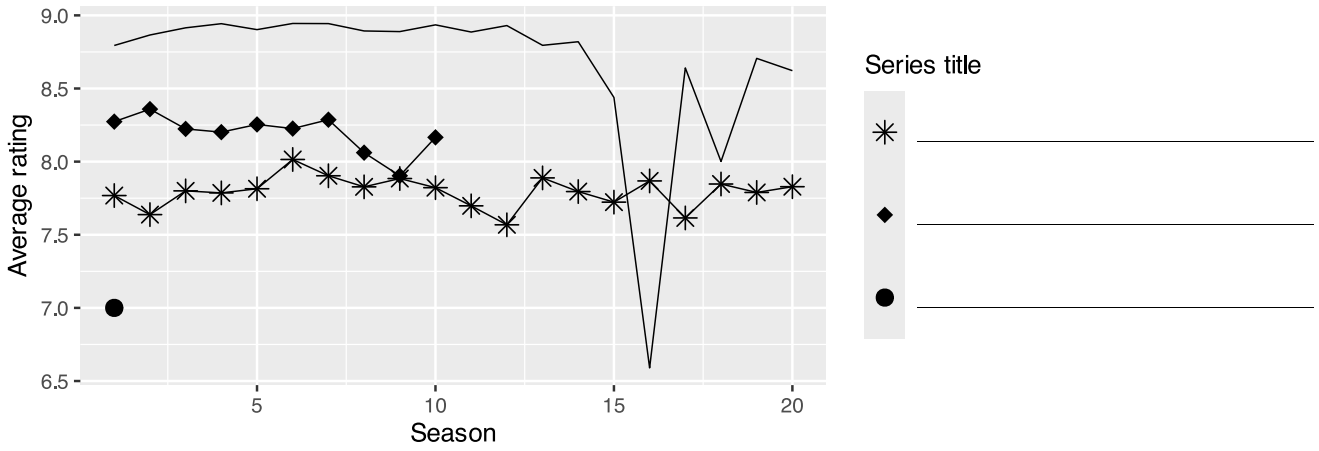
You will work with data on average ratings for each season of three series from the Law & Order-verse – a subset of the data from the previous questions. Below is a peek at the first ten rows of the Law & Order data.

The plot below shows the distributions of average ratings of various Law & Order series across seasons.



Question 23

Based on the information from the side-by-side box plots, fill in the legend of the plot below with Law & Order series titles.



[1] "Doink doink" is the scene and episode introductory sound on the Law & Order series. If you've never heard it, you're not at any disadvantage for the exam. If you've ever heard it, good luck getting it out of your head!

Question 24

The following code calculates the standard deviations of average season ratings of the five Law & Order series. Unfortunately, the output is partially erased and replaced with blanks.

```
lo_titles <- c(
  "Law & Order",
  "Law & Order: Criminal Intent",
  "Law & Order: SVU"
)

law_and_order |>
  filter(title %in% lo_titles) |>
  group_by(title) |>
  summarize(mean_av_rating = mean(av_rating), sd_av_rating = sd(av_rating))
```

```
# A tibble: 5 × 3
  title                          mean_av_rating sd_av_rating
  <chr>                                   <dbl>        <dbl>
1 Law & Order                            _(1)_        0.106
2 Law & Order: Criminal Intent            8.20        0.129
4 Law & Order: SVU                        8.67        _(2)_
```

Based on the visualizations you've seen of these data so far, which of the following is **true** about the blanks in the output? **Select all that are true.**

a. The **mean** of average ratings (Blank 1) of Law & Order seasons is **lower** than the other two means.

b. The **mean** of average ratings (Blank 1) of Law & Order seasons is **higher** than the other two means.

c. The **standard deviation** of average ratings of Law & Order: SVU seasons (Blank 2) is **lower** than the other two standard deviations.

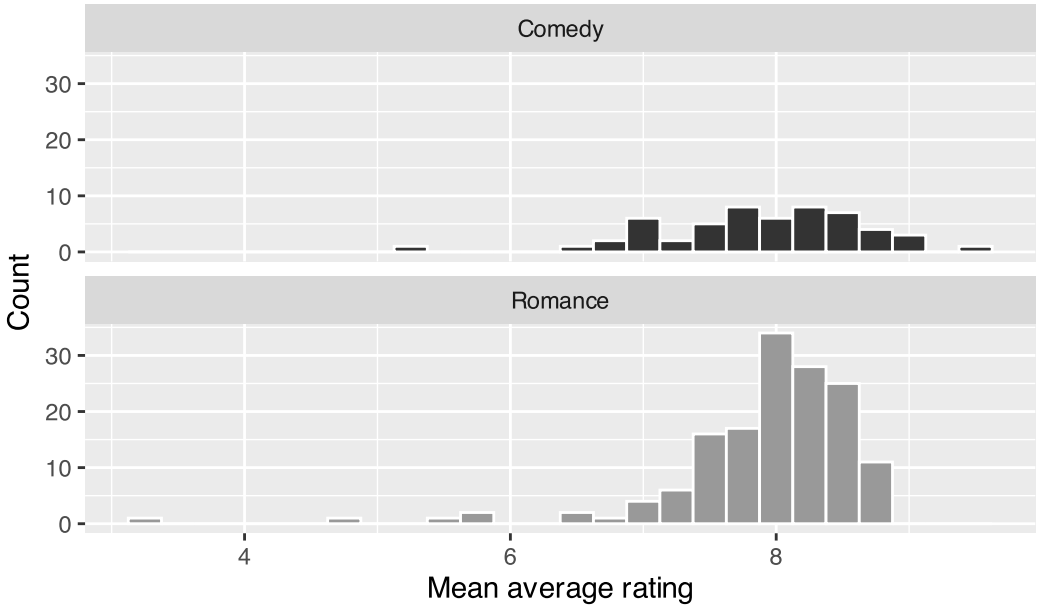
d. The **standard deviation** of average ratings of Law & Order: SVU seasons (Blank 2) is **higher** than the other two standard deviations.

e. The **standard deviation** of average ratings of Law & Order: SVU seasons (Blank 2) is **between** the other two standard deviations.

Romance and comedy

Finally, we focus on romance and comedy shows. We first filter the dataset for any shows that have romance or comedy as their genre (`genre_1`, `genre_2`, or `genre_3`) and then remove shows that have both of these genre labels. For the next two questions, we focus on these shows that we identify as either romance or comedy. We then calculate the mean of the average season ratings for each show, to obtain a single "mean average rating" value per show.

The plot below shows the distributions of mean average ratings of seasons of comedy and romance shows.



Question 25

Which of the following statements is **true** about these distributions? **Select all that are true.**

a. Mean average ratings of romance shows are bimodal.
b. Mean average ratings of comedy are unimodal.
c. Mean average ratings of romance shows is left skewed.
d. Mean average ratings of comedy shows is right skewed.
e. There are more romance shows than comedy shows.

IMDB

The data for the next few questions come from the Internet Movie Database (IMDB). Specifically, the data are a random sample of movies released between 1980 and 2020.

The name of the data frame used for this analysis is `movies`, and it contains the variables shown in Table 1.

Table 1: Data dictionary for `movies`

Variable	Description
`name`	name of the movie
`rating`	rating of the movie (R, PG, etc.)
`genre`	main genre of the movie.
`runtime`	duration of the movie
`year`	year of release
`release_date`	release date (YYYY-MM-DD)
`release_country`	release country
`score`	IMDB user rating
`votes`	number of user votes
`director`	the director
`writer`	writer of the movie
`star`	main actor/actress
`country`	country of origin
`budget`	the budget of a movie (some movies don't have this, so it appears as 0)
`gross`	revenue of the movie
`company`	the production company

The first thirty rows of the `movies` data frame are shown in Table 2, with variable types suppressed (since we'll ask about them later).

Table 2
First 30 rows of `movies`, with variable types suppressed.

A tibble: 500 × 16

	name	score	runtime	genre	rating	release_country	release_date	budget	gross	votes	year	director	writer	star	company	country
1	Blue City	4.4	83 mins	Action	R	United States	1986-05-02	10000000	6947787	1100	1986	Michelle Man…	Ross Macdona…	Judd Nelson	Paramount Pi…	United States
2	Winter Sleep	8.1	196	Drama	Not Rated	Turkey	2014-06-12	NA	4018705	48000	2014	Nuri Bilge C…	Ebru Ceylan	Haluk Bilgin…	NBC Film	Turkey
3	Rang De Basanti	8.1	167	Comedy	Not Rated	United States	2006-01-26	NA	10800778	115000	2006	Rakeysh Ompr…	Renzil D'Sil…	Aamir Khan	ROMP	India
4	Pokémon Detecti…	6.6	104	Action	PG	United States	2019-05-10	150000000	433921300	146000	2019	Rob Letterman	Dan Hernandez	Ryan Reynolds	Warner Bros.	United States
5	A Bad Moms Chri…	5.6	104	Comedy	R	United States	2017-11-01	28000000	130560428	46000	2017	Jon Lucas	Jon Lucas	Mila Kunis	Huayi Brothe…	United States
6	Replicas	5.5	107	Drama	PG-13	United States	2019-01-11	30000000	9330075	34000	2018	Jeffrey Nach…	Chad St. John	Keanu Reeves	Company Films	United Kingd…
7	Windy City	5.8	103	Drama	R	Uruguay	1986-01-01	NA	343890	262	1984	Armyan Berns…	Armyan Berns…	John Shea	CBS Theatric…	United States
8	War for the Pla…	7.4	140	Action	PG-13	United States	2017-07-14	150000000	490719763	235000	2017	Matt Reeves	Mark Bomback	Andy Serkis	Twentieth Ce…	United States
9	Tales from the …	6.4	98	Crime	R	United States	1995-05-24	6000000	11837928	7400	1995	Rusty Cundie…	Rusty Cundie…	Clarence Wil…	40 Acres & A…	United States
10	Fire with Fire	6.5	103	Drama	PG-13	United States	1986-05-09	NA	4636169	1500	1986	Duncan Gibbi…	Bill Phillips	Craig Sheffer	Paramount Pi…	United States
11	Raising Helen	6	119	Comedy	PG-13	United States	2004-05-28	50000000	49718611	36000	2004	Garry Marsha…	Patrick J. C…	Kate Hudson	Touchstone P…	United States
12	Feeling Minneso…	5.4	99	Comedy	R	United States	1996-09-13	NA	3124440	11000	1996	Steven Baige…	Steven Baige…	Keanu Reeves	Fine Line Fe…	United States
13	The Babe	5.9	115	Biography	PG	United States	1992-04-17	NA	19930973	9300	1992	Arthur Hiller	John Fusco	John Goodman	Universal Pi…	United States
14	The Real Blonde	6	105	Comedy	R	United States	1998-02-27	NA	83488	3900	1997	Tom DiCillo	Tom DiCillo	Matthew Modi…	Lakeshore En…	United States
15	To vlemma tou O…	7.6	176	Drama	Not Rated	United States	1997-11-01	NA	NA	6400	1995	Theodoros An…	Theodoros An…	Harvey Keitel	Greek Film C…	Greece
16	Going the Dista…	6.3	102	Comedy	R	United States	2010-09-03	32000000	42059111	57000	2010	Nanette Burs…	Geoff LaTuli…	Drew Barrymo…	New Line Cin…	United States
17	Jung on zo	6.8	103	Action	R	Hong Kong	1993-06-24	NA	3741869	6100	1993	Kirk Wong	Tin Nam Chun	Jackie Chan	Golden Harve…	Hong Kong
18	Rita, Sue and B…	6.5	93	Comedy	R	United Kingdom	1987-05-29	NA	124167	3600	1987	Alan Clarke	Andrea Dunbar	Siobhan Finn…	British Scre…	United Kingd…
19	Phone Booth	7	81	Crime	R	United States	2003-04-04	13000000	97837138	255000	2002	Joel Schumac…	Larry Cohen	Colin Farrell	Fox 2000 Pic…	United States
20	Happy Death Day	6.6	96	Comedy	PG-13	United States	2017-10-13	4800000	125479266	124000	2017	Christopher …	Scott Lobdell	Jessica Rothe	Universal Pi…	United States
21	Barely Legal	4.7	90	Comedy	R	Thailand	2006-05-25	NA	83439	5900	2003	David Mickey…	David H. Ste…	Erik von Det…	Internationa…	Germany
22	Three Kings	7.1	114	Action	R	United States	1999-10-01	75000000	107752036	163000	1999	David O. Rus…	John Ridley	George Cloon…	Warner Bros.	United States
23	Menace II Socie…	7.5	97	Crime	R	United States	1993-05-26	3500000	27912072	54000	1993	Albert Hughes	Allen Hughes	Tyrin Turner	New Line Cin…	United States
24	Four Rooms	6.8	98	Comedy	R	United States	1995-12-25	4000000	4257354	100000	1995	Directors	Allison Ande…	Tim Roth	Miramax	United States
25	Quartet	6.8	98	Comedy	PG-13	United States	2013-03-01	11000000	59520298	19000	2012	Dustin Hoffm…	Ronald Harwo…	Maggie Smith	Headline Pic…	United Kingd…
26	Tape	7.2	86	Drama	R	Denmark	2002-07-12	100000	515900	19000	2001	Richard Link…	Stephen Belb…	Ethan Hawke	Detour Filmp…	United States
27	Marked for Death	6	93	Action	R	United States	1990-10-05	12000000	57968936	21000	1990	Dwight H. Li…	Michael Grais	Steven Seagal	Victor & Gra…	United States
28	Congo	5.2	109	Action	PG-13	United States	1995-06-09	50000000	152022101	43000	1995	Frank Marsha…	Michael Cric…	Laura Linney	Paramount Pi…	United States
29	Stop-Loss	6.4	112	Drama	R	United States	2008-03-28	25000000	11212953	20000	2008	Kimberly Pei…	Mark Richard	Ryan Phillip…	Paramount Pi…	United States
30	Con Air	6.9	115	Action	R	United States	1997-06-06	75000000	224012234	282000	1997	Simon West	Scott Rosenb…	Nicolas Cage	Touchstone P…	United States

ℹ 470 more rows

Question 26

The `name` and `runtime` variables are shown below, with the variable types suppressed.

```
# A tibble: 500 × 2
  name                      runtime
1 Blue City                 83 mins
2 Winter Sleep              196
3 Rang De Basanti           167
4 Pokémon Detective Pikachu 104
5 A Bad Moms Christmas      104
6 Replicas                  107
# ℹ 494 more rows
```

What is the type of the `runtime` variable?

a. Character

b. Double

c. Factor

d. Integer

e. Logical

Question 27

The code below summarizes the data in a certain way.

```
movies |>
  summarize(sum(release_country == "United States"))
```

```
# A tibble: 1 × 1
  `sum(release_country == "United States")`
                                      <int>
1                                       435
```

Which of the following is **TRUE** about the code and its result? **Select all that are true.**

a. Evaluates whether each `release_country` is equal to `"United States"` or not, which results in a logical variable.

b. Filters out rows where `release_country` is not equal to `"United States"` and counts the remaining rows.

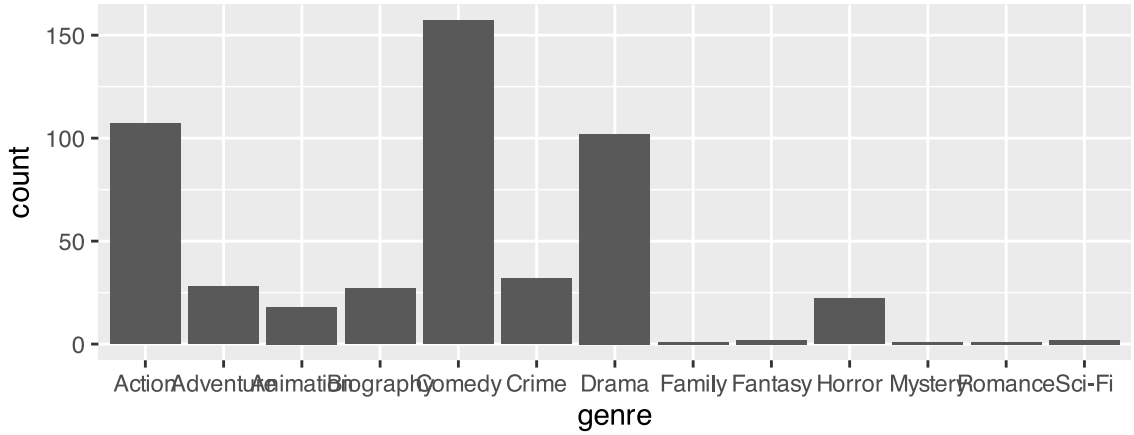
c. Sums the logical values, where each `TRUE` is considered a 1 and each `FALSE` is considered a `0`.

d. Results in a character vector.

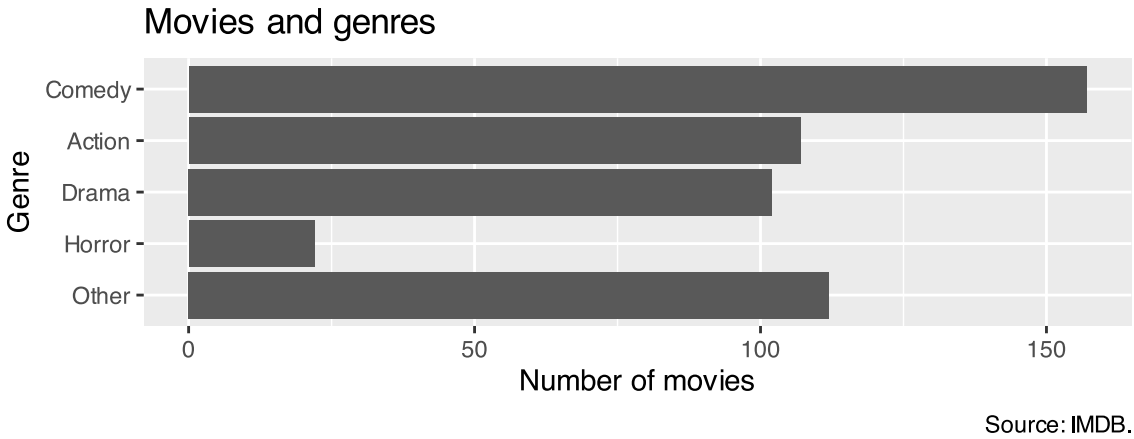
e. The result shows there are 435 movies released in the United States.

Question 28

Suppose you want a visualization that shows the number of movies in the sample in each `genre`. Your first attempt is as follows.



A friend of yours says that the visualization is difficult to read and they suggest using the following visualization instead.



Which of the following modifications would your friend have made to your code to create their version? **Select all that apply.**

a. Combine movies in genres other than Comedy, Drama, Action, and Horror into a new level called `"Other"`.

b. Reorder the levels in descending order of numbers of observations, except for the `"Other"` level.

c. Map `genre` to the `y` aesthetic.

d. Add a title, x and y-axis labels, and a caption.

e. Filter out all moves in genres other than Comedy, Drama, Action, and Horror before plotting.

Question 29

Which of the following is **TRUE** about the code and its result? **Select all that are true.**

```
movies |>
  count(rating, genre) |>
  pivot_wider(names_from = genre, values_from = n, values_fill = 0)
```

```
# A tibble: 6 × 6
  rating    Other Drama Action Comedy Horror
  <fct>     <int> <int>  <int>  <int>  <int>
1 G             5     1      1      1      0
2 PG           38    13     10     18      0
3 PG-13        19    25     35     35      0
4 R            45    50     57     96     21
5 NC-17         1     2      0      1      0
6 Not Rated     4    11      4      6      1
```

a. The code counts how many movies are in each rating and genre combination.

b. The code sorts the results in descending order.

c. Each row of the output is a movie.

d. The output shows that there are six distinct ratings in the dataset.

e. The code reduces the number of variables and observations in the `movies` data frame to six.

Gymnastics

The `gymnastics` dataset contains information about routines performed by gymnasts in the women's artistic gymnastics competition at the 2024 Paris Olympics. Each row represents a single routine, so a gymnast who competed on all four apparatus appears in at least four rows. Variables include the gymnast's `name`, `country`, `event` (apparatus), `competition` (round), and `score` (routine total).

```
Rows: 561
Columns: 5
$ name        <chr> "Simone Biles", "Simone Biles", "Rebeca Andrade", "Rebeca …
$ country     <chr> "United States", "United States", "Brazil", "Brazil", "Uni…
$ event       <chr> "Vault", "Vault", "Vault", "Vault", "Vault", "Vault", "Vau…
$ competition <chr> "Qualification", "Qualification", "Qualification", "Qualif…
$ score       <dbl> 15.800, 14.800, 14.900, 14.466, 14.666, 14.200, 14.333, 14…
```

The `continents` dataset contains countries (`entity`) and the `continent` each one is in.

```
Rows: 285
Columns: 2
$ entity    <chr> "Abkhazia", "Afghanistan", "Akrotiri and Dhekelia", "Aland I…
$ continent <chr> "Asia", "Asia", "Asia", "Europe", "Europe", "Africa", "Ocean…
```

A student joins the two datasets and summarizes average vault score by continent with the following pipeline:

```
gymnastics |>
  left_join(continents, join_by(country == entity)) |>
  filter(event == "Vault") |>
  group_by(continent) |>
  summarize(mean_vault_score = mean(score, na.rm = TRUE)) |>
  arrange(desc(mean_vault_score))
```

```
# A tibble: 7 × 2
  continent     mean_vault_score
  <chr>                    <dbl>
1 South America             14.1
2 North America             14.1
3 <NA>                      13.7
4 Europe                    13.5
5 Asia                      13.5
6 Oceania                   13.4
7 Africa                    13.3
```

Question 30

Part 1: The resulting summary has a row where `continent` is `NA`. Explain what this row represents — i.e., what is happening in the join that produces it?

Part 2: The student then runs the following code and gets the output shown:

```
gymnastics |>
  anti_join(continents, join_by(country == entity)) |>
  distinct(country)
```

```
# A tibble: 3 × 1
  country
  <chr>
1 Great Britain
2 AUT
3 Chinese Taipei
```

What is the `mean_vault_score` of athletes from the Great Britain, Austria, and Chinese Taipei? Provide a justification for how you arrived at this number.

Inflation around the world

The `country_inflation` dataset contains annual inflation rates for countries around the world from 1993 to 2025. The first five rows and a handful of the columns are shown below:

```
# A tibble: 5 × 6
  country   `1993` `2000` `2010` `2020` `2025`
  <chr>      <dbl>  <dbl>  <dbl>  <dbl>  <dbl>
1 Australia   1.75   4.46   2.92  0.847   2.87
2 Austria     3.63   2.34   1.81  1.38    3.53
3 Belgium     2.75   2.54   2.19  0.741   2.47
4 Canada      1.87   2.72   1.78  0.717   2.07
5 Chile      12.7    3.84   1.41  3.05   NA
```

A student reshapes the data with the following code:

```
country_inflation_long <- country_inflation |>
  pivot_longer(
    cols = !country,
    names_to = "year",
    names_transform = as.numeric,
    values_to = "inflation_rate"
  )
```

and then uses `country_inflation_long` to make the following plot:

```
Warning: Removed 1 row containing missing values or values outside the scale range
(`geom_line()`).
```

```
Warning: Removed 1 row containing missing values or values outside the scale range
(`geom_point()`).
```

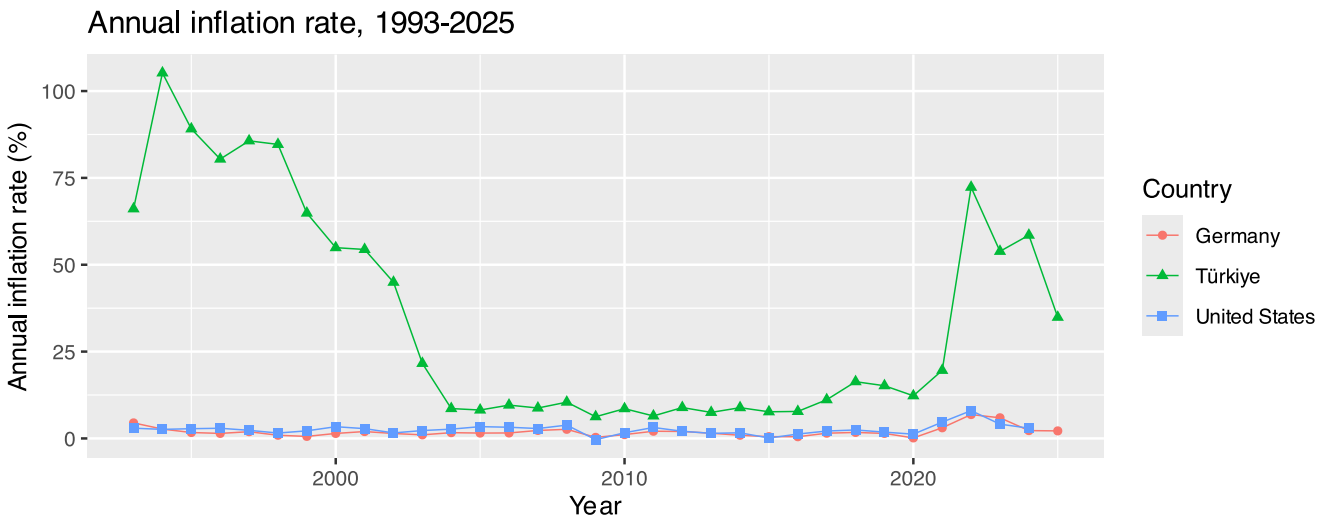


Figure 4: Annual inflation rate for selected countries, 1993-2025

Question 31

The original `country_inflation` dataset has 38 rows and 34 columns. How would you calculate the number of rows and columns `country_inflation_long` has, and what does each row represent? Explain, in your own words, what `pivot_longer()` did to the data, including what the `names_transform = as.numeric` argument accomplished.

Question 32

A student wants to compare each country's inflation rate in 2025 to its rate in 1993 (the first and last years in the `country_inflation` dataset). They write the following pipeline:

```
country_inflation |>
  filter(!is.na(`1993`), !is.na(`2025`)) |>
  mutate(inf_ratio = `2025` / `1993`) |>
  select(country, inf_ratio) |>
  arrange(desc(inf_ratio)) |>
  slice_head(n = 5)
```

```
# A tibble: 5 × 2
  country        inf_ratio
  <chr>              <dbl>
1 New Zealand         2.20
2 Australia           1.64
3 Denmark             1.51
4 Ireland             1.50
5 United Kingdom      1.48
```

Part 1: In 1-2 sentences, describe what the resulting data frame represents and how to interpret the values in `inf_ratio` (e.g., what does an `inf_ratio` of 2.20 mean for New Zealand?).

Part 2: A teammate says the `filter()` step is unnecessary because "`mutate()` works fine with missing values." Explain what would be different about the resulting data frame if the `filter()` step were removed, and why that could be misleading.

Part 3: Another teammate suggests replacing `mutate()` with `summarize()` in this pipeline, keeping everything else the same. Explain why that would not accomplish the student's goal.

Part 4: That teammate is dead set on using `summarize()` instead of `mutate()`. What other modification would they need to make to the code to get the same result?

Part 5: The student then removes `slice_head(n = 5)` from their pipeline so the result contains all countries, saves it as `inflation_ratios`, and makes the following plot:

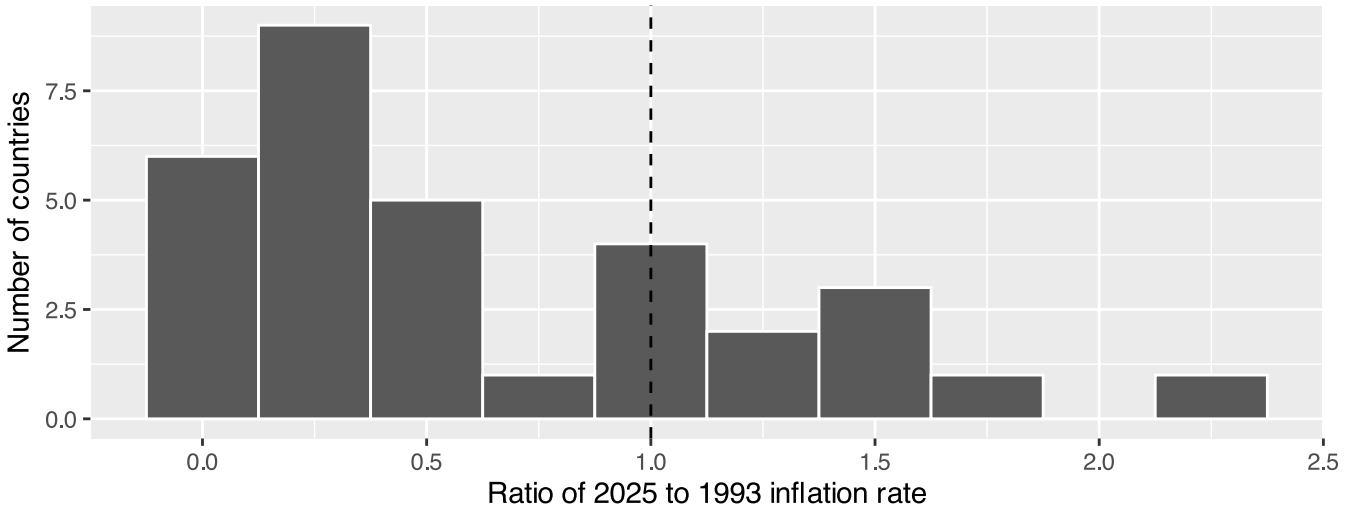


Figure 5: Ratio of 2025 to 1993 inflation rates

In a few sentences, describe the distribution shown in Figure 5, explain what the dashed vertical line at 1 represents, and state what it means for the countries to the left vs. the right of that line.

Bonus

Pick a concept we introduced in class so far that you've been struggling with and explain it in your own words.